



elastic

A report on :

“ ELastic stack integration with Honeypot - Tpot “

“ Hamza Jameel “
(dev.hamzaj@gmail.com)

Table of contents :

Introduction:	3
Installations and Configurations :	3
1. Honeypot - Tpot:	3
2. ELK Stack:	5
HoneyPot Integration with ElasticSearch:	7
1. Pre requisites:	7
• A fleet server on ELK up and running	7
• A windows and Linux Fleet policies	7
• An ELK agent installed on the Tpot installed endpoint.	8
2. Data Filtration on ELK:	8
3. Visualization:	9
Summary:	9

Introduction:

In this project, I successfully implemented a monitoring and detection setup using the Elastic Stack (ELK). This solution revolves around detecting OS ticket queries, streamlines incident tracking and management through features like customizable workflows, automated ticket routing, and real-time notifications. By integrating osTicket with Elastic SIEM, the security of organizations can elevate , enabling seamless incident detection, tracking, and resolution. This integration allows security alerts generated by Elastic SIEM to be automatically converted into actionable tickets in osTicket, ensuring prompt response and effective collaboration.

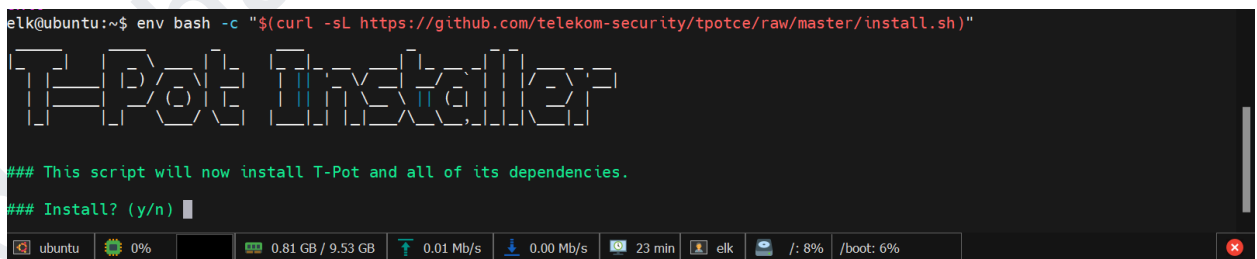
Installations and Configurations :

This setup requires two different installations of different stacks . In this lab, the installations are performed on the Ubuntu base system.

1. Honeypot - Tpot:

T-Pot Honeypot is an open-source, multi-layered honeypot solution designed to capture and analyze malicious activities on a network. It combines multiple honeypot technologies into a single platform, offering a wide range of decoy services to attract and monitor attackers. T-Pot is built on Docker containers and includes tools like **Cowrie**, **Dionaea**, **Elastic Stack (ELK)**, and **Suricata** to collect detailed logs and alerts from various attack vectors. To move forward , install the tpot with single command installation [script](#) :

```
env bash -c "$(curl  
-sLhttps://github.com/telekom-security/tpotce/raw/master/install.sh)"
```



```
elk@ubuntu:~$ env bash -c "$(curl -sL https://github.com/telekom-security/tpotce/raw/master/install.sh)"
T-Pot Installer
## This script will now install T-Pot and all of its dependencies.
## Install? (y/n) █
ubuntu 0% 0.81 GB / 9.53 GB 0.01 Mb/s 0.00 Mb/s 23 min elk /: 8% /boot: 6%
```

After this , the installation script will start running and user need to enter some credentials for this , remember this installation script is running on the top of docker engine .

```
## Creating base64 encoded htpasswd username and password for T-Pot config file: /home/elk/tpotce/.env

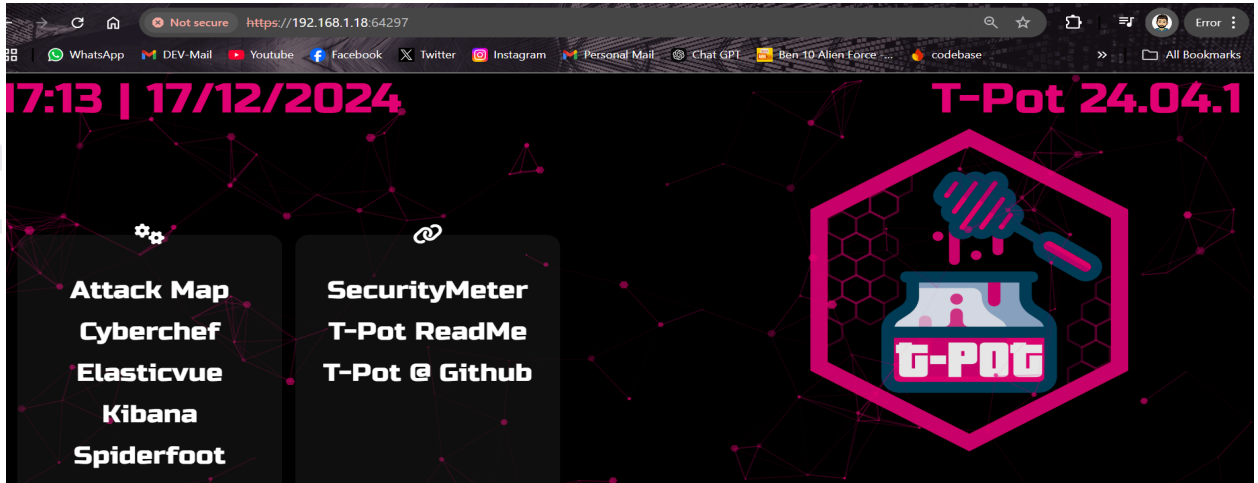
## Now pulling images ...
[+] Pulling 14/26
✓ conpot_ipmi Skipped - Image is already being pulled by conpot_guardian_ast 0.0s
✓ map_data Skipped - Image is already being pulled by map_web 0.0s
✓ conpot_kamstrup_382 Skipped - Image is already being pulled by conpot_guardian_ast 0.0s
✓ conpot_IEC104 Skipped - Image is already being pulled by conpot_guardian_ast 0.0s
: dicompot [ ] Pulling 4.9s
: conpot_guardian_ast [ ] Pulling 4.9s
: suricata [ ] Pulling 4.9s
✓ medpot Pulled 3.6s
: logstash [##] 189.1MB / 500.5MB Pulling 4.9s
: elasticsearch [ ] Pulling 4.9s
: p0f [ ] Pulling 4.9s
: honeytrap [ ] Pulling 4.9s
: spiderfoot [ ] Pulling 4.9s
: ciscoasa [ ] Pulling 4.9s
: map_web [.] 16.35MB / 39.98MB Pulling 4.9s
: kibana [ ] Pulling 4.8s
✓ nginx Pulled 4.1s
: ewsposter [ ] Pulling 4.8s
: tpotinit [ ] Pulling 4.8s
: honeypots [ ] Pulling 4.8s
: map_redis [ ] Pulling 4.8s
: fati [ ] Pulling 4.8s
: adbhoney [ ] Pulling 4.8s
```

The tpot is installed and running now . Navigate to the Ubuntu command line and type the netstat command for open ports .

```
Active Internet connections (only servers)
Proto Recv-Q Send-Q Local Address           Foreign Address         State       User        Inode         PID/Program name
tcp        0      0 127.0.0.1:6010          0.0.0.0:*                LISTEN      1000        12923        1072/sshd: elk@pts/
tcp        0      0 127.0.0.1:6791          0.0.0.0:*                LISTEN      0           1635         723/elastic-agent
tcp        0      0 127.0.0.1:6789          0.0.0.0:*                LISTEN      0           7821         723/elastic-agent
tcp        0      0 0.0.0.0:64295           0.0.0.0:*                LISTEN      0           49289        14797/sshd: /usr/sb
tcp6       0      0 :::6010                  :::*                    LISTEN      1000        12922        1072/sshd: elk@pts/
tcp6       0      0 :::64295                 :::*                    LISTEN      0           49291        14797/sshd: /usr/sb

## Done. Please reboot and re-connect via SSH on tcp/64295.
```

Note : **SSH** port has been moved from port **22** to port **64295**. It is now necessary to connect from this port to access the system.



Although the dashboards are already provided in Tpot, the custom installation is recommended for hands-on lab and custom parsing .

2. ELK Stack:

Download the elastic search latest file from the official repository or from their webpage . In my case I am using the base operating system linux ubuntu. For ELK setup , two different modules needed to be downloaded: the elastic search and kibana . Then they both will be interlinked to share and visualize the data . Use the wget command to download the package and then use dpkg -i command to install the elastic search into the target system.

```
root@BAA0FR-ELK:~# dpkg --get-selections | grep elasticsearch
Preparing to unpack elasticsearch-8.15.0-amd64.deb ...
Creating elasticsearch group... OK
Creating elasticsearch user... OK
Unpacking elasticsearch (8.15.0) ...
Setting up elasticsearch (8.15.0) ...
-----*----- Security autoconfiguration information -----
Authentication and authorization are enabled.
TLS for the transport and HTTP layers is enabled and configured.

The generated password for the elastic built-in superuser is : cju3j08XZNLHUzoKZ7rr

If this node should join an existing cluster, you can reconfigure this with
'/usr/share/elasticsearch/bin/elasticsearch-reconfigure-node --enrollment-token <token-here>'
after creating an enrollment token on your existing cluster.

You can complete the following actions at any time:

Reset the password of the elastic built-in superuser with
'/usr/share/elasticsearch/bin/elasticsearch-reset-password -u elastic'.

Generate an enrollment token for Kibana instances with
'/usr/share/elasticsearch/bin/elasticsearch-create-enrollment-token -s kibana'.

Generate an enrollment token for Elasticsearch nodes with
'/usr/share/elasticsearch/bin/elasticsearch-create-enrollment-token -s node'.

#####
### NOT starting on installation, please execute the following statements to configure elasticsearch service to start automatically using systemd
sudo systemctl daemon-reload
```

After successful installation it will show some useful credentials that we will need to set up our kibana instance. Note these credentials and save them in a notepad file. After successful setup of our elastic instance we will setup kibana instance.

Download the kibana debian file from the official webpage and use the command dpkg -i to unpack it.

After that , move towards the elastic search directory and create an enrollment token for kibana as shown:

HoneyPot Integration with ElasticSearch:

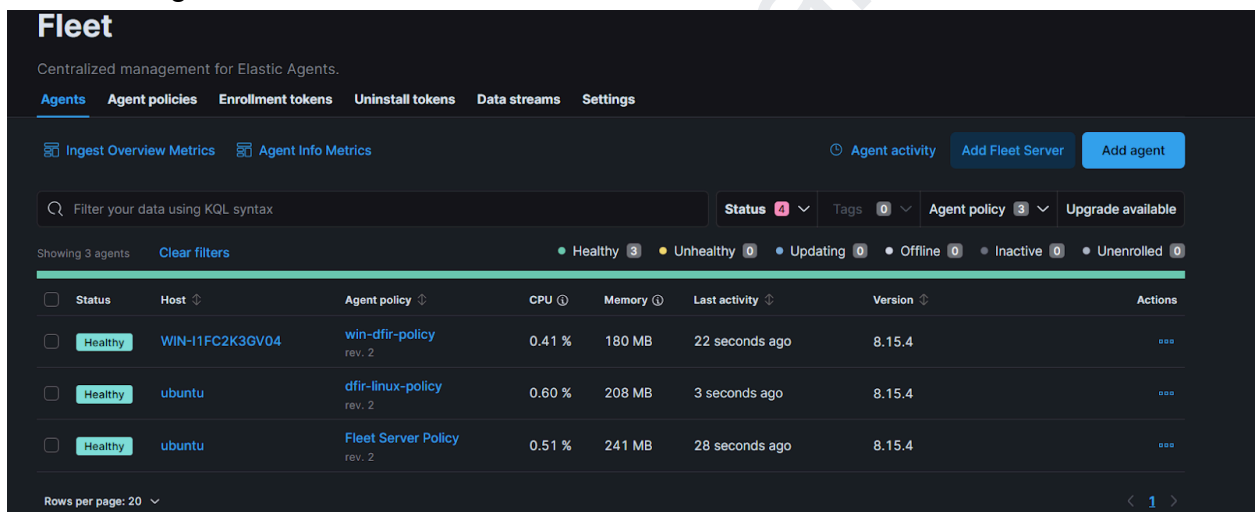
To parse the data generated by mythic server into elastic stack you need the following things :

1. Pre requisites:

Some of the components which are required before the practical are:

- **A fleet server on ELK up and running**

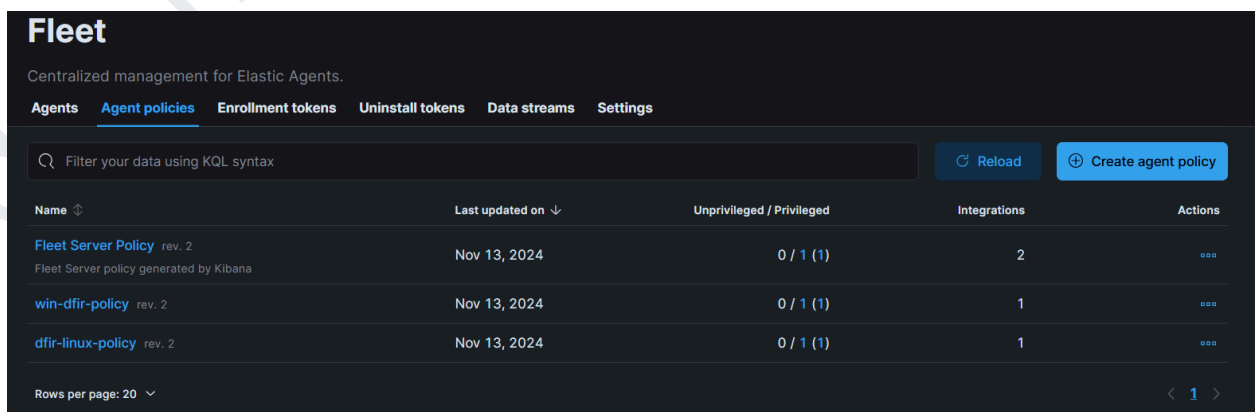
Navigate to your Kibana instance and then click on 3 lines on the top left corner. Navigate to Fleet at last and set up the fleet policies. From the agent option click on install agent at any endpoint and then also add the flag `--insecure` in case you are not using the ssl certificates.



The screenshot shows the Kibana Fleet interface. The top navigation bar includes 'Agents', 'Agent policies', 'Enrollment tokens', 'Uninstall tokens', 'Data streams', and 'Settings'. The 'Agents' tab is selected. Below the navigation bar, there are buttons for 'Ingest Overview Metrics', 'Agent Info Metrics', 'Agent activity', 'Add Fleet Server', and 'Add agent'. A search bar is present with the text 'Filter your data using KQL syntax'. Below the search bar, there are filters for 'Status' (4), 'Tags' (0), 'Agent policy' (3), and 'Upgrade available'. A summary bar shows 'Showing 3 agents' and a 'Clear filters' button. Below this, there is a table with columns: Status, Host, Agent policy, CPU, Memory, Last activity, Version, and Actions. The table lists three agents: 'WIN-11FC2K3GV04' (Healthy), 'ubuntu' (Healthy), and 'ubuntu' (Healthy). The bottom of the page shows 'Rows per page: 20' and a pagination control.

Status	Host	Agent policy	CPU	Memory	Last activity	Version	Actions
Healthy	WIN-11FC2K3GV04	win-dfir-policy rev. 2	0.41 %	180 MB	22 seconds ago	8.15.4	...
Healthy	ubuntu	dfir-linux-policy rev. 2	0.60 %	208 MB	3 seconds ago	8.15.4	...
Healthy	ubuntu	Fleet Server Policy rev. 2	0.51 %	241 MB	28 seconds ago	8.15.4	...

- **A windows and Linux Fleet policies**



The screenshot shows the Kibana Fleet interface with the 'Agent policies' tab selected. The top navigation bar is the same as the previous screenshot. Below the navigation bar, there are buttons for 'Reload' and 'Create agent policy'. A search bar is present with the text 'Filter your data using KQL syntax'. Below the search bar, there is a table with columns: Name, Last updated on, Unprivileged / Privileged, Integrations, and Actions. The table lists three policies: 'Fleet Server Policy rev. 2', 'win-dfir-policy rev. 2', and 'dfir-linux-policy rev. 2'. The bottom of the page shows 'Rows per page: 20' and a pagination control.

Name	Last updated on	Unprivileged / Privileged	Integrations	Actions
Fleet Server Policy rev. 2 Fleet Server policy generated by Kibana	Nov 13, 2024	0 / 1 (1)	2	...
win-dfir-policy rev. 2	Nov 13, 2024	0 / 1 (1)	1	...
dfir-linux-policy rev. 2	Nov 13, 2024	0 / 1 (1)	1	...

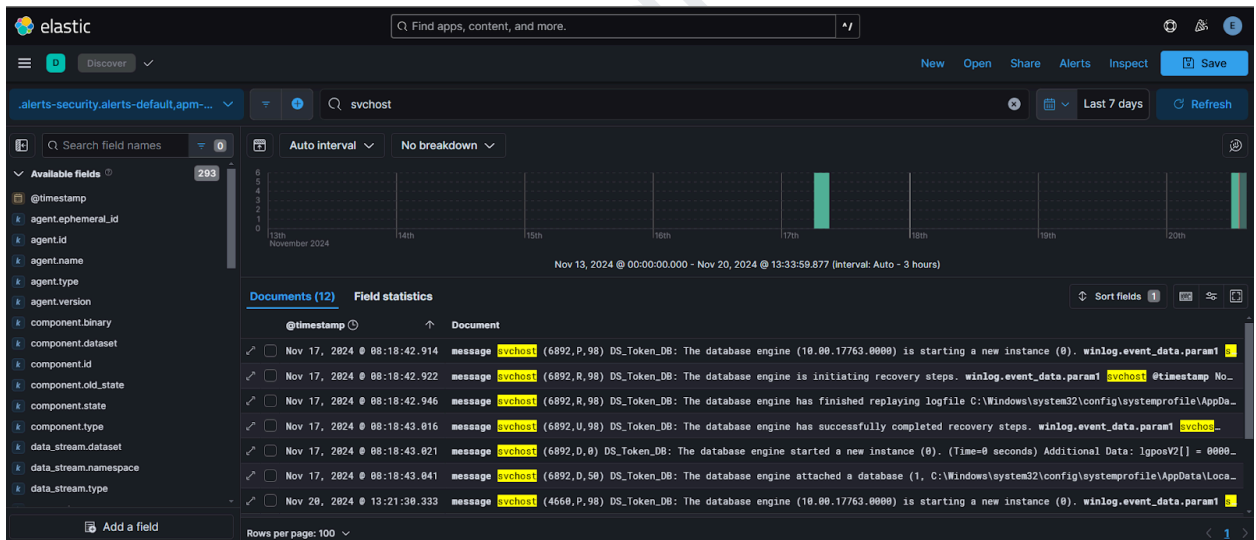
- An ELK agent installed on the Tpot installed endpoint.

Status	Host	Agent policy	CPU	Memory	Last activity	Version	Actions
Healthy	WIN-11FC2K3GV04	win-dfir-policy rev. 2	0.42 %	180 MB	3 seconds ago	8.15.4	...

This will allow you to parse the syslogs from honeypot os endpoint to the ELK deployed stack. After ensuring these instances , you can proceed with the next steps.

2. Data Filtration on ELK:

Once an agent is installed on the endpoint , we can cross check whether the data from the endpoint is coming to the kibana interface or not. For this , select the index pattern is discover tab and use the keywords of the service we have installed on the endpoint.

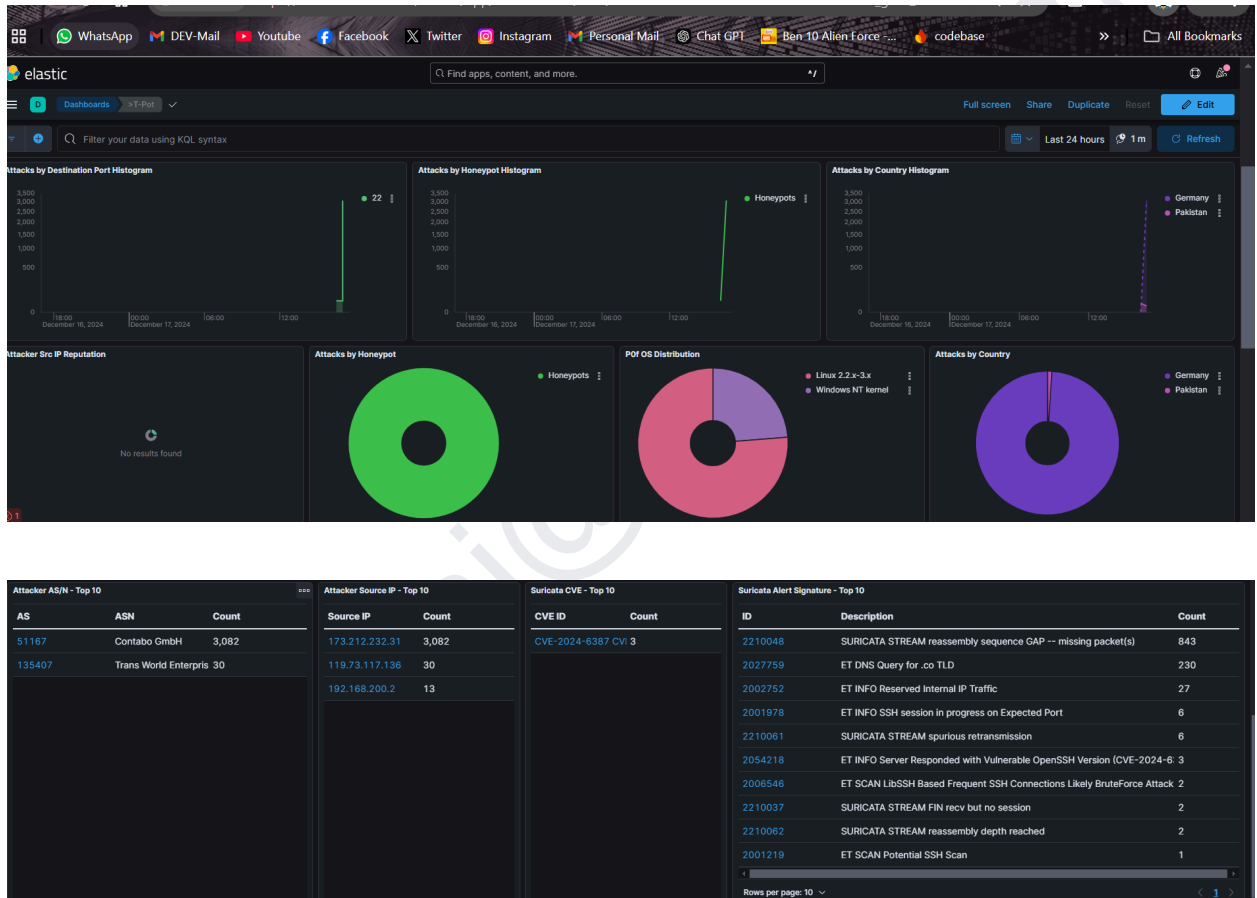


As visible in the screenshot the data is coming and parsing is enabled . Now we can apply different queries to filter the required data . After applying some queries and filtering the data you will get the data in the discover tab . Save the data with queries and now the custom dashboards on use cases can be made

Important note : Although the dashboards are already provided in Tpot but the custom installation is recommended for hands-on lab and custom parsing .

3. Visualization:

The live events and attacks can be visualized at kibana like this :



Furthermore, this project can be extended to capture the IOC's like users can add their own threat intelligence databases and can use python scripts to train the machine learning models for future use .

Summary:

The project involved setting up a Fleet Server on ELK for centralized agent management, configuring a honeypot server to simulate adversarial techniques, and installing an Elastic Agent on Windows endpoints for log collection. Using ELK's detection rules and alerting mechanisms.

dev.hamzaj@gmail.com